



Project Documentation

PoC 64-bit Operating System — ELF + WinPE Support (MIT Licensed)

Languages: x86_64 Assembly & C

Test Platform: QEMU

Bootloader: GRUB2 (external)

Filesystem: Custom EXT4 driver + mkfs

Binary Support: Custom ELF and WinPE Parsers

License: MIT



Part 1: Legal Documentation

1.1 Project Licensing and Identity

- **License:** MIT
- **Project Status:** Proof of Concept (PoC)
- **Development:** Entire kernel, binary loaders, EXT4 FS, and tools are original.
- **Third-Party Bootloader:** GRUB2 used externally for multiboot2-compliant loading.

MIT License Declaration:

MIT License

Copyright (c) 2024-2025 COS

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

1.2 Third-Party Usage and Scope

GRUB2 only loads the kernel; the kernel and drivers are not GPL-contaminated. EXT4 support (driver + mkfs) is original and not derived from Linux kernel code. QEMU is used for testing only, not redistributed or integrated. No reuse of Windows or Linux code in ELF/PE parsers.

Component	Source/Origin	License Type of Use Notes
GRUB2 Bootloader	GNU GRUB	GPLv3
QEMU Emulator	QEMU Project	GPLv2
EXT4 Specification	Kernel Docs	Public Domain Functional reference
PE/ELF Specs	MS / SysV ABI	Public Docs
	Public Docs	
	(Intel Manual, OSDev) and	
ASM Boot Chain	Open Source Projects	Public Domain Documentation of Several Aspects of the Boot Sequence
	Explanations	

It should be noted that a C reimplementation has been considered for the final release due to the similarities present within the source code and other projects as the boot sequence may only be made following steps that are always the same.

1.3 IP Risk & Compliance Mitigation

Binary formats:

- ELF and PE are supported through specs and observation, not code reuse.
- Windows-specific behavior is avoided unless publicly documented.

Filesystem:

- EXT4 support based on structure, not code lifting.

Bootloader:

- GRUB2 is treated as an external loader; avoid bundling to preserve MIT purity.



PART 2: TECHNOLOGY WATCH (2025)

Ongoing awareness of tools, formats, kernel development, and licensing updates.

2.1 Monitored Topics

Date	Topic	Source	Summary
2025-05-24	Recursive Mapping	https://os.phil-opp.com/page-tables/#recursive-mapping	Clever trick where the page tables map themselves — a local fix that simplifies page-table manipulation. This was added during the creation of the memory mapping algorithm as we realised that we needed a way to access every virtual physical address from the nodes forming the page table.
2025-04-14	GCC Stack Canaries	https://ctf-wiki.mahaloz.re/pwn/linux/mitigation/canary/	Explains how GCC inserts secret values into the stack to

Date	Topic	Source	Summary
			catch and halt buffer overflows. During compilation, we encountered an issue where GCC would add __chk_fail() around some of the functions. An incomplete implementation needed to be provided.
2025-02-13	RTC Initialisation for EXT4 Metadata	https://web.archive.org/web/20150514082645/http://www.nondot.org/sabre/os/files/MiscHW/RealtimeClockFAQ.txt	Details how to work with the Real-Time Clock. We decided to implement proper support for the RTC to provide correct timestamps for the EXT4 filesystem.
2025-01-11	Understanding Virtual Paging	https://zolulat.github.io/understanding-paging/	Explains how the system translates virtual addresses into physical ones.
2024-11-30	Memory Management	https://wiki.osdev.org/Memory_Map_(x86)	Lays out the physical memory map on x86. It became relevant once the MemoryManagement module could finally be addressed properly.
2024-11-17	Disk ATA PIO Mode (48bit)	https://wiki.osdev.org/ATA_PIO_Mode	A guide to issuing low-level disk commands via ATA PIO using 48-bit addressing. Our team had to write a driver for ATA disks to be able to communicate with them.

Date	Topic	Source	Summary
2024-11-17	ATA Registers	https://atola.com/products/tf1/manual/registers.html	with Qemu's virtual hardware.
2024-11-10	Higher Half Setup	https://medium.com/@connorstack/how-does-a-higher-half-kernel-work-107194e46a64	Inventory of I/O ports and registers needed to communicate directly with ATA drives.
2024-10-06	OS Initialisation	https://www.gnu.org/software/grub/manual/multiboot/multiboot.html	Explains why and how kernels are mapped high in virtual memory — separating them from userspace cleanly. This was implemented as a preventative measure long before the Memory Manager was actually started.
2024-09-15	EXT4 Specification	https://www.kernel.org/doc/html/v6.9/filesystems/ext4/index.html?highlight=ext4	The Multiboot specification — the handshake between bootloaders and kernels, ensuring smooth entry into the land. This structure needed to be properly written to ensure that GRUB would return proper hardware information, including the free and reserved memory block addresses.
			Technical exposition of EXT4's internal mechanisms, as

Date	Topic	Source	Summary
2024-09-15	EXT4 Disk Layout	https://blogs.oracle.com/linux/post/understanding-ext4-disk-layout-part-1	implemented in the Linux kernel. Walks through how EXT4 organizes data physically on disk. It was used during the initial implementation of the <code>cos_mkfs</code> function as well as the read/write syscalls.
2024-09-08	Translation Lookaside Buffer	https://en.wikipedia.org/wiki/Translation_lookaside_buffer	The TLB is a cache for memory mappings that speeds up address resolution. TLB flush sequences were added long before the Memory Manager was completed as a preventative measure once we would reach virtual address mapping.
2024-09-02	Allowed Freestanding Headers	https://wiki.osdev.org/C_Library#Freestanding_and_Hosted	Lists the standard library headers permissible when one's compiling without an OS — freestanding environment. This link was especially useful as we were not sure whether or not we could use the <code>va_args</code> functionality usually provided by complete OSes.
2024-08-31			Deep dive into a legacy interrupt

Date	Topic	Source	Summary
	8285A PIC Documentation	https://pdos.csail.mit.edu/6.828/2012/readings/hardware/8259A.pdf	controller. It was especially relevant considering the fact that we had trouble initialising it at first.
2024-08-31	8295A PIC Initialisation Procedure	http://www.brokenthorn.com/Resources/OSDev16.html	Walkthrough for initializing the PIC. It sets the stage for interrupt-driven execution. This step was required as enabling long mode without it would invariably end up triggering Triple Faults due to the hardware timer on IRQ0 being improperly remapped.
2024-08-31	AMD64 Red Zone	https://forum.osdev.org/viewtopic.php?t=21720	The AMD64 Red Zone is a quiet 128-byte region below the stack pointer that shouldn't be touched in leaf functions - per the AMD64 ABI. We recall that during the initialisation process, we had a hint of an issue setting the PIC and researched into the matter leading to the discovery of several compile flaws. We should have integrated into the project on the kernel side at the very least.
2024-08-31	Better Debugger	https://sourceforge.net/projects/bochs/files/bochs/2.8/	Bochs includes a detailed software debugger tailored

Date	Topic	Source	Summary
2024-07-04 and 2024-08-31	IDT Documentation	https://wiki.osdev.org/Interrupt_Descriptor_Table	emulating and inspecting low-level x86 internals. It proved useful for debugging hardware related issues we encountered.
2024-07-04 and 2024-08-31	GDT Documentation	https://wiki.osdev.org/Global_Descriptor_Table	Covers how to write up the table that routes hardware software interrupts to handlers. It was required for initialisation of the hardware.
2024-06-24	ASM Instructions	https://www.felixcloutier.com/x86/	Guide to building Global Descriptor Table. It was required for initialisation of the hardware.
2024-06-24	OS Initialisation	https://cdrdv2.intel.com/v1/dl/getContent/671447	A rich catalog of x86 instruction sets — invaluable for decoding binary behavior and crafting assembly. Intel's roadmap from power-on to OS handoff — BIOS, firmware, and early hardware setup.

2.2 Future considered enhancements/features

Date	Topic	Source	Summary
2025-05-11	Zero Page	https://en.wikipedia.org/wiki/Zero_page	At some point, we asked ourselves how complete OSes handled NULL

Date	Topic	Source	Summary
2025-05-04	Program Launching	https://electronics.stackexchange.com/questions/440421/int-functions-with-no-return-statement	<p>pointer dereferences as Linux defined it as (void *)0. We ended up discovering that this Virtual Address acted as a trap instead and should be reserved by the kernel. It still hasn't been implemented within COS.</p> <p>Once program launching became a concern once more, we asked ourselves how complete OSes would handle main() functions with no return, as the CPU would end up reading garbage if it was called as a function pointer. It is then that we realised that launching a full program</p>

Date	Topic	Source	Summary
2025-04-08	Memory Management Optimisation/ Feature	https://www.kernel.org/doc/gorman/html/understand/understand011.html	wouldn't be as simple as using function pointers to call the main() function and instead that we would need a scheduler to gracefully handle context switching. Compared to our custom Memory Manager, this model would pre-partition the RAM into caches managing slabs (blocks of contiguous memory). With these caches, select commonly used objects can be reused without the need to reallocate them first, therefore speeding up operations. For instance, allocating a structure and then asking the kernel to

Date	Topic	Source	Summary
2025-02-13	ELF Specification	https://refspecs.linuxfoundation.org/elf/elf.pdf	<p>free it would instead keep the structure in memory [with its internal information being reinitialised] only if it is commonly accessed.</p> <p>Blueprint for the ELF binary format — the standard for Unix-like systems' executables and object files. This will be required for the next part of the project and thus hadn't been implemented yet.</p>
2024-09-15	Hash Tree Algorithm	https://blogs.oracle.com/linux/post/understanding-ext4-disk-layout-part-2	<p>This blog post discusses the Hash Trees used by the EXT4 specification for storing directory entries. By default, they are linearly written in directory data</p>

Date	Topic	Source	Summary
2024-09-08	Memory Management Optimisation/ Feature	https://en.wikipedia.org/wiki/Copy-on-write	<p>blocks. This newer way of indexing files and folders would allow for the search complexity to be logarithmic instead.</p> <p>Technique that delays memory duplication until modification — efficient for process creation and memory savings. Once scheduling tasks will be implemented, shared libraries (which are read-only resources) could be accessed by several programs in a memory efficient way.</p>
2024-09-08	Memory Management Optimisation	https://www.scaler.com/topics/operating-system/page-replacement-algorithm/	<p>A tour of strategies for deciding which memory pages get evicted — central to managing</p>

Date	Topic	Source	Summary
2024-08-31	APIC Documentation	https://web.archive.org/web/20140308064246/http://www.osdever.net/tutorials/pdf/apic.pdf	RAM pressure. This Memory Management technique has been considered as an evolution of the project where multitasking is available. Intro to the Advanced Programmable Interrupt Controller — key for multiprocessor and fine-grained interrupt control. This is the newest hardware used as part of Intel CPUs, replacing the PIC. We would like to update the code with proper APIC support later down the line.
2024-06-21	WinPE File Format Specification	https://learn.microsoft.com/en-us/windows/win32/debug/pe-format	Microsoft's breakdown of the PE binary format. During the planning of project, we started work on a WinPE binary parser.

2.3 Excluded Enhancements/Features

Date	Topic	Source	Summary
2024-09-06	Level 5 Paging	https://www.intel.com/content/www/us/en/content-details/671442/5-level-paging-and-5-level-ept-white-paper.html	Discusses the expansion of virtual address space with 5-level paging — useful when dealing with colossal memory ranges (more than 256TiB). Due to the fact that the COS will never need that much memory to function properly, this technology was discarded early on during the planning phase of the project.

🤝 PART 3: EXPERT INTERACTIONS & CONTACT LOG

3.1 Expert Directory

Name / Alias	Specialty	Affiliation / Project	Contact	Status
Dr. Moy	Compilation	Independent	Matthieu.Moy@univ-lyon1.fr	Active

3.2 Conferences & Community Engagement

Contacts/discussion boards which have been considered but haven't yet been reached out for: - OSDev Forums (Help related to OS Development and low level) - CrossOver Team (Main contributor to the WINE project, a translation layer for Linux)

APPENDICES

Yearly Development Timeline

Start date	Target date	Status
2025-05-31	2025-06-05	Mickael: Adding legal documentation. Corentin + Mickael: Continuing work on MemoryManager. Thomas + Adam: Continuing remaining work on FileSystem.
2025-05-13	2025-05-30	Mickael + Corentin: Fixing Memory Mapping and adding proper Kernel Mapping (to be tested thoroughly). Adam: Continuing work on touch. Thomas: Continuing work on mv.
2025-04-29	2025-05-13	Mickael + Corentin: Adding Memory Allocation and starting work on Memory Mapping. Thomas: Continuing work on mv. Adam: Continuing work on touch.
2025-04-14	2025-04-29	Mickael: Adding ls command PoC. Corentin + Mickael: Continuing Memory

Start date	Target date	Status
2025-03-31	2025-04-14	<p>Management Revised algorithm (Page Frame Manager and Reclaiming Allocated Unused Kernel Pages)</p> <p>Adam: Starting work on the touch command.</p> <p>Thomas: Starting work on the mv command.</p> <p>Adam + Thomas: Adding full support for EXT4 read capabilities.</p> <p>Adam + Thomas: Adding full read/write wrappers.</p> <p>Mickael: Fixing the sectorcount bug + Merging of every branches into one.</p> <p>Corentin: Starting the revision of the MemoryManagement algorithm (Page Frame Manager and Free Memory Claiming)</p> <p>Mickael: Adding support for stack smashing canaries being added by GCC.</p> <p>Fixing the memset bug by adding a properly named memset function.</p>
2025-03-15	2025-03-27	<p>Adam + Thomas: Continuing work on the read/write wrappers.</p>
2025-02-28	2025-02-15	

Start date	Target date	Status
2025-02-11	2025-02-24	<p>Adam: Fixing a part of the write wrapper + Working on read wrapper.</p> <p>Thomas: Working on write wrapper.</p> <p>Corentin: Adding more printf arguments and updating the command line to support arguments.</p> <p>Mickael: Fixing the EXT4 Support.</p> <p>Adam + Thomas: Starting work on the read and write wrappers respectfully.</p>
2025-01-31	2025-02-13	<p>Adam + Corentin: Fixed Block Descriptor entries and Block Descriptor Bitmaps.</p> <p>Mickael: Added Inode Headers and Inode Bitmaps. Initialisation of default inodes</p> <p>Added support for RTC to the COS in order for timestamps to work.</p> <p>Thomas: Started research on Linux ELF parsing.</p>
2025-01-01	2025-01-11	<p>Thomas: Fixed the read/write ASM functions.</p> <p>Mickael: Fixed the read/write ASM functions + Added</p>

Start date	Target date	Status
2024-12-14	2024-12-15	<p>incomplete Virtual Addressing support.</p> <p>Corentin: Work on MemoryManagement.</p> <p>Adam: Work on Ext4 filesystem.</p> <p>Thomas: Debugging of the Read/Write function phase.</p> <p>Mickael: Started working on Documenting and sourcing the project under the Documentation module.</p>
2024-12-07	2024-12-08	<p>Corentin: Work on MemoryManagement.</p> <p>Adam: Work on Ext4 filesystem.</p> <p>Thomas: Debugging of the Read/Write function phase.</p> <p>Mickael: Started working on Documenting and sourcing the project under the Documentation module.</p>
2024-11-30	2024-12-01	<p>Corentin: Started work on MemoryManagement by creating necessary structures.</p> <p>Adam: Continued work on Ext4 filesystem.</p> <p>Thomas: Debugging of the Read/Write function phase.</p> <p>Mickael: Started work</p>

Start date	Target date	Status
2024-11-23	2024-11-24	<p>on RAM block ordering in memory.</p> <p>Corentin: Tried to fix the mirroring to push every module to the EPITECH mirror.</p> <p>Adam: Continued work on Ext4 filesystem.</p> <p>Thomas: Added Read/Write ASM functions to the core module.</p> <p>Testing phase.</p> <p>Mickael: Tried to fix the mirroring to push every module to the EPITECH mirror + Helped Thomas on debugging.</p> <p>Corentin: Work on Unit Testing support and some code refactoring work.</p> <p>Adam: Continued work on Ext4 Filesystem.</p>
2024-11-16	2024-11-17	<p>Thomas: Continued work on Read/Write ASM functions.</p> <p>Mickael: Finished working on Higher Half Kernel + RAM space retrieval thanks to GRUB.</p> <p>Corentin: Work on Unit Testing support and initial push to the EPITECH mirror.</p>
2024-11-09	2024-11-10	<p>Adam: Continued work on Ext4 Filesystem blocks.</p>

Start date	Target date	Status
2024-11-02	2024-11-03	<p>Thomas: Continued work on Read/Write ASM functions.</p> <p>Mickael: Helped Adam on Ext4 filesystem blocks + continued work on Higher Half Kernel.</p> <p>Corentin: Worked on adding the workflow to the project.</p> <p>Adam: Continued work on Ext4 Filesystem blocks.</p> <p>Thomas: Continued work on Read/Write ASM functions.</p> <p>Mickael: Sending the Kernal to Higher Half Memory.</p>